

# LINEAR ALGEBRA

## The Complete Beginner's Guide

From Vectors to Machine Learning

Step-by-Step • Practice Problems • ML Applications

&#128218; 10 Chapters &#9999; 50+ Practice Problems &#129302; ML Connections &#128161; Visual Explanations

Prepared for Joel • February 2026

# Table of Contents

---

<b>Chapter 1</b>	<b>Scalars, Vectors &amp; Their Operations</b>
	The building blocks
<b>Chapter 2</b>	<b>Vector Spaces &amp; Linear Combinations</b>
	Span, basis, and independence
<b>Chapter 3</b>	<b>Matrices: The Core Tool</b>
	Creation, types, and basic operations
<b>Chapter 4</b>	<b>Matrix Multiplication</b>
	The most important operation
<b>Chapter 5</b>	<b>Linear Transformations</b>
	What matrices really do
<b>Chapter 6</b>	<b>Systems of Linear Equations</b>
	Solving $Ax = b$
<b>Chapter 7</b>	<b>Determinants</b>
	Measuring scale and invertibility
<b>Chapter 8</b>	<b>Eigenvalues &amp; Eigenvectors</b>
	The hidden structure
<b>Chapter 9</b>	<b>Norms, Dot Products &amp; Distances</b>
	Measuring in vector spaces
<b>Chapter 10</b>	<b>Putting It All Together: ML Applications</b>
	PCA, regression, and neural nets

# How to Use This Guide

---

This guide is designed to take you from zero knowledge to a solid understanding of linear algebra, with a special focus on the concepts most relevant to machine learning. Here's what each section contains:

■ **Key Insight** - Important concepts and intuitions you should remember

❖ **Practice Problems** - Hands-on exercises to test your understanding

■ **Answers** - Solutions to practice problems - try solving first!

■ **ML Connection** - How this concept is used in machine learning

**Difficulty Progression:** Each chapter builds on the previous one. Chapters 1-3 are **Easy**, Chapters 4-6 are **Medium**, and Chapters 7-10 are **Advanced**. Take your time with each level before moving on.

## Chapter 1

# Scalars, Vectors & Their Operations

EASY

## 1.1 What is a Scalar?

A **scalar** is simply a single number. That's it! Examples: 5, -3.14, 0.001, 100. In math, we usually write scalars as lowercase italic letters like  $a$ ,  $b$ ,  $c$ .

### Definition: Scalar

A scalar is a single numerical value. It represents a magnitude (size) without direction. In programming, think of it as a single variable holding one number:  $x = 5$

## 1.2 What is a Vector?

A **vector** is an ordered list of numbers. Think of it as a container that holds multiple numbers in a specific order. Each number in the vector is called a **component** or **element**.

**Physical intuition:** A vector represents both a **direction** and a **magnitude** (length). Imagine an arrow pointing from one place to another - the direction it points and how long it is together define the vector.

2D Vector	3D Vector	n-D Vector
$v = [3, -5]$ x-component: 3 y-component: -5	$v = [3, -5, 4]$ x: 3, y: -5, z: 4	$v = [v_1, v_2, \dots, v_n]$ n components

### ■ Key Insight

In machine learning, data is stored as vectors. A single data point with 10 features becomes a 10-dimensional vector. An image with 784 pixels becomes a 784-dimensional vector. The math works the same regardless of dimension!

## 1.3 Vector Addition

---

To add two vectors, simply add their corresponding components. Both vectors **must have the same dimension** (same number of components).

$$[3, -5] + [2, 1] = [3+2, -5+1] = [5, -4]$$

$$[3, -5, 4] + [2, 1, -3] = [3+2, -5+1, 4+(-3)] = [5, -4, 1]$$

**Physical intuition:** Vector addition is like following directions. Walk along vector A, then walk along vector B. Where you end up is the result of  $A + B$ . This is called the "tip-to-tail" method.

#### Properties of vector addition:

- **Commutative:**  $a + b = b + a$  (order doesn't matter)
- **Associative:**  $(a + b) + c = a + (b + c)$  (grouping doesn't matter)
- **Zero vector:**  $a + 0 = a$  (adding zero changes nothing)

## 1.4 Scalar Multiplication

To multiply a vector by a scalar, multiply **every component** by that scalar.

$$2 * [3, -5] = [2*3, 2*(-5)] = [6, -10]$$

$$-1 * [3, -5] = [-3, 5] \text{ (reverses direction!)}$$

$$0.5 * [3, -5] = [1.5, -2.5] \text{ (halves the length)}$$

**Physical intuition:** Scalar multiplication **stretches** or **shrinks** a vector. Multiply by 2 = twice as long. Multiply by 0.5 = half as long. Multiply by -1 = same length but opposite direction.

## 1.5 Vector Subtraction

Subtraction is just addition with a negated vector:

$$a - b = a + (-1 * b)$$

$$[3, -5] - [2, 1] = [3-2, -5-1] = [1, -6]$$

## ■ Machine Learning Connection

In machine learning, these basic operations are everywhere:

- **Feature vectors:** Each data point is a vector. A house described by [price, sqft, bedrooms] = [250000, 1500, 3]
- **Weight updates:** In gradient descent, weights are updated by:  $w = w - \text{learning\_rate} * \text{gradient}$  (scalar multiplication + subtraction!)
- **Batch processing:** Adding vectors lets you combine predictions from multiple models

```
Python: import numpy as np  
v1 = np.array([3, -5])  
v2 = np.array([2, 1])  
print(v1 + v2) # [5, -4]  
print(2 * v1) # [6, -10]
```

## ☞ Practice Problems

**P1.1** Calculate:  $[4, -2, 7] + [1, 5, -3]$

**P1.2** Calculate:  $3 * [2, -1, 4]$

**P1.3** Calculate:  $[10, 5] - [3, 8]$

**P1.4** If  $v = [1, 2, 3]$  and  $w = [4, 5, 6]$ , compute  $2v + 3w$

**P1.5** A robot moves  $[3, 2]$  then  $[-1, 4]$  then  $[2, -3]$ . Where does it end up?

## ■ Answers

**A1.1**  $[5, 3, 4]$

**A1.2**  $[6, -3, 12]$

**A1.3**  $[7, -3]$

**A1.4**  $2*[1,2,3] + 3*[4,5,6] = [2,4,6] + [12,15,18] = [14, 19, 24]$

**A1.5**  $[3+(-1)+2, 2+4+(-3)] = [4, 3]$

## Chapter 2

# Vector Spaces & Linear Combinations

EASY

## 2.1 Linear Combinations

A **linear combination** takes multiple vectors, scales each by a scalar, and adds them together. This is the most fundamental operation in all of linear algebra.

### Definition: Linear Combination

Given vectors  $v_1, v_2, \dots, v_n$  and scalars  $c_1, c_2, \dots, c_n$ , the linear combination is:  $c_1v_1 + c_2v_2 + \dots + c_nv_n$

**Example:** Given  $v_1 = [1, 0]$  and  $v_2 = [0, 1]$ :

$$3*[1, 0] + 2*[0, 1] = [3, 0] + [0, 2] = [3, 2]$$

The scalars 3 and 2 are called **coefficients** or **weights**.

## 2.2 Span

The **span** of a set of vectors is the collection of ALL possible linear combinations of those vectors. It answers the question: "What points can I reach using these vectors?"

### ■ Key Insight

**Think of it like mixing paint colors:**

- With just red paint, you can only make shades of red (a line)
- With red AND blue paint, you can make reds, blues, and purples (a plane)
- With red, blue, AND yellow, you can reach almost any color (3D space)

Each new independent vector "opens up" a new dimension of possibility.

## 2.3 Linear Independence

Vectors are **linearly independent** if none of them can be written as a linear combination of the others. In other words, each vector adds a genuinely new direction.

Independent ■	Dependent ■
[1,0] and [0,1] <i>Neither is a multiple of the other. They point in completely different directions.</i>	[1,2] and [2,4] <i>[2,4] = 2*[1,2]. The second vector is just the first stretched by 2. Same direction!</i>

## 2.4 Basis

A **basis** is a set of linearly independent vectors that spans the entire space. It's the minimum set of vectors needed to describe every point in the space.

The **standard basis** for 2D is:  $e_1 = [1, 0]$  and  $e_2 = [0, 1]$

The **standard basis** for 3D is:  $e_1 = [1, 0, 0]$ ,  $e_2 = [0, 1, 0]$ ,  $e_3 = [0, 0, 1]$

Any 2D vector can be written using the standard basis:  $[3, -5] = 3*[1,0] + (-5)*[0,1]$

## 2.5 Dimension

The **dimension** of a vector space is the number of vectors in any basis. 2D space has dimension 2 (need 2 basis vectors), 3D space has dimension 3, and so on.

### ■ Machine Learning Connection

Linear combinations are the foundation of machine learning:

- **Neural network output:** Each neuron computes a linear combination of inputs:  $y = w_1x_1 + w_2x_2 + \dots + w_nx_n + \text{bias}$
- **Feature engineering:** Creating new features from existing ones is often a linear combination
- **Linear regression:** The prediction is literally a linear combination:  $y = w_1 * \text{feature}_1 + w_2 * \text{feature}_2 + \dots$
- **Word embeddings:** king - man + woman ≈ queen (vector arithmetic on word vectors!)

### Practice Problems

- P2.1** Express  $[7, 11]$  as a linear combination of  $[1, 2]$  and  $[3, 1]$
- P2.2** Are  $[1, 2, 3]$  and  $[2, 4, 6]$  linearly independent? Why or why not?
- P2.3** Can you reach the point  $[5, 5]$  using only multiples of  $[1, 1]$ ? What about  $[3, 7]$ ?
- P2.4** How many vectors are in the standard basis for 5D space?

### Answers

- A2.1** Solve:  $c_1*[1,2] + c_2*[3,1] = [7,11]$ . From row 1:  $c_1 + 3c_2 = 7$ . From row 2:  $2c_1 + c_2 = 11$ . Solving:  $c_1 = 4$ ,  $c_2 = 1$ . Check:  $4*[1,2] + 1*[3,1] = [4,8] + [3,1] = [7,11]$  ■
- A2.2** No!  $[2,4,6] = 2*[1,2,3]$ . They are linearly dependent (same direction, different length).
- A2.3**  $[5,5] = 5*[1,1]$  ■. But  $[3,7]$  cannot be reached - any multiple of  $[1,1]$  has equal components.
- A2.4** 5 vectors (one for each dimension).

## Chapter 3

# Matrices: The Core Tool

EASY

## 3.1 What is a Matrix?

A **matrix** is a rectangular grid of numbers arranged in rows and columns. A matrix with  $m$  rows and  $n$  columns is called an  **$m \times n$  matrix** (read as " $m$  by  $n$ ").

### Definition: Matrix

An  $m \times n$  matrix  $A$  has  $m$  rows and  $n$  columns. We write  $A_{ij}$  to refer to the element in row  $i$ , column  $j$ . Matrices are usually written with capital bold letters: **A, B, C**.

2x3 Matrix	3x2 Matrix	2x2 Matrix
$\begin{vmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{vmatrix}$ <i>2 rows, 3 columns</i>	$\begin{vmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{vmatrix}$ <i>3 rows, 2 columns</i>	$\begin{vmatrix} 1 & 2 \\ 3 & 4 \end{vmatrix}$ <i>Square matrix</i>

## 3.2 Special Types of Matrices

Type	Description	Example
Identity ( $I$ )	1s on diagonal, 0s elsewhere. Like the number 1 for matrices.	$\begin{vmatrix} 1 & 0 \\ 0 & 1 \end{vmatrix}$
Zero Matrix	All elements are 0. Like the number 0.	$\begin{vmatrix} 0 & 0 \\ 0 & 0 \end{vmatrix}$
Diagonal	Non-zero only on main diagonal.	$\begin{vmatrix} 3 & 0 \\ 0 & 7 \end{vmatrix}$
Symmetric	$A = A^T$ (same when flipped over diagonal).	$\begin{vmatrix} 1 & 2 \\ 2 & 3 \end{vmatrix}$

Transpose ( $A^T$ )	Rows become columns, columns become rows.	$\begin{vmatrix} 1 & 2 \\ 3 & 4 \end{vmatrix}^T = \begin{vmatrix} 1 & 3 \\ 2 & 4 \end{vmatrix}$
---------------------	---	---

### 3.3 Matrix Addition & Scalar Multiplication

Just like vectors! Add element-by-element (matrices must be same size), or multiply every element by a scalar.

$$\begin{array}{rcl} \begin{vmatrix} 1 & 2 \\ 3 & 4 \end{vmatrix} + \begin{vmatrix} 5 & 6 \\ 7 & 8 \end{vmatrix} & = & \begin{vmatrix} 1+5 & 2+6 \\ 3+7 & 4+8 \end{vmatrix} = \begin{vmatrix} 6 & 8 \\ 10 & 12 \end{vmatrix} \\[1em] 3 * \begin{vmatrix} 1 & 2 \\ 3 & 4 \end{vmatrix} & = & \begin{vmatrix} 3 & 6 \\ 9 & 12 \end{vmatrix} \end{array}$$

#### ■ Key Insight

A vector is actually a special case of a matrix! A column vector with  $n$  elements is an  $n \times 1$  matrix. A row vector is a  $1 \times n$  matrix. This is why the same rules apply to both.

#### ■ Machine Learning Connection

Matrices are the language of machine learning:

- **Datasets:** Your entire dataset is a matrix! Each row is a data point, each column is a feature.
- **Weight matrices:** Neural network layers store their weights in matrices.
- **Images:** A grayscale image is a matrix of pixel values (0-255). A color image is 3 matrices (R, G, B).

```
Python: dataset = np.array([[5.1, 3.5, 1.4], # iris sample 1
                           [4.9, 3.0, 1.4], # iris sample 2
                           [4.7, 3.2, 1.3]]) # iris sample 3
```

### Practice Problems

**P3.1** What is the size of a matrix with 4 rows and 3 columns?

**P3.2** Add:  $|1\ 2\ 3| + |4\ 5\ 6|$ . Write each row.

$$|4\ 5\ 6| \quad |7\ 8\ 9|$$

**P3.3** Find the transpose of:  $|1\ 2\ 3|$

$$|4\ 5\ 6|$$

**P3.4** Is  $|1\ 3|$  symmetric? Why or why not?

$$|3\ 5|$$

**P3.5** If your dataset has 1000 samples and 50 features, what size is the data matrix?

### ■ Answers

**A3.1** 4 x 3 (4-by-3)

**A3.2**  $|5\ 7\ 9| / |11\ 13\ 15|$

**A3.3**  $|1\ 4| / |2\ 5| / |3\ 6|$  (3x2 matrix - rows and columns swap)

**A3.4** Yes! Flipping over the diagonal gives the same matrix ( $A=A^T$ ).

**A3.5** 1000 x 50

## Chapter 4

# Matrix Multiplication

MEDIUM

## 4.1 The Key Rule: Dot Product

Before matrix multiplication, you need to understand the **dot product** of two vectors. Multiply corresponding elements and add them up:

$$[1, 2, 3] \cdot [4, 5, 6] = 1*4 + 2*5 + 3*6 = 4 + 10 + 18 = 32$$

The dot product takes two vectors and returns a **single number** (a scalar).

## 4.2 How Matrix Multiplication Works

Matrix multiplication is NOT element-by-element! Instead, each element of the result is a **dot product** between **a row of the first matrix and a column of the second**.

### ■ Key Insight

**The Size Rule:** To multiply A ( $m \times n$ ) by B ( $n \times p$ ), the number of columns in A MUST equal the number of rows in B. The result will be an  $m \times p$  matrix.

- (2x3) times (3x4) = (2x4) ■ The 3s match!
- (2x3) times (2x4) = ERROR ■ 3 does not equal 2!

Step-by-step example:

$$\begin{vmatrix} 1 & 2 \end{vmatrix} \times \begin{vmatrix} 5 & 6 \end{vmatrix} = ?$$
$$\begin{vmatrix} 3 & 4 \end{vmatrix} \quad \begin{vmatrix} 7 & 8 \end{vmatrix}$$

Position	Calculation	Result
Row 1, Col 1	$[1,2] \cdot [5,7] = 1*5 + 2*7$	19
Row 1, Col 2	$[1,2] \cdot [6,8] = 1*6 + 2*8$	22
Row 2, Col 1	$[3,4] \cdot [5,7] = 3*5 + 4*7$	43

Row 2, Col 2

$$[3,4] \cdot [6,8] = 3*6 + 4*8$$

50

$$\text{Result} = \begin{vmatrix} 19 & 22 \end{vmatrix}$$

$$\begin{vmatrix} 43 & 50 \end{vmatrix}$$

## 4.3 Important Properties

- **NOT commutative:**  $AB$  does NOT equal  $BA$  in general! Order matters!
- **Associative:**  $(AB)C = A(BC)$
- **Distributive:**  $A(B + C) = AB + AC$
- **Identity:**  $AI = IA = A$  (the identity matrix is like multiplying by 1)

## 4.4 Matrix-Vector Multiplication

A very common operation: multiplying a matrix by a vector. This is just matrix multiplication where the second matrix has only one column.

$$\begin{vmatrix} 2 & 1 \end{vmatrix} \times \begin{vmatrix} 3 \end{vmatrix} = \begin{vmatrix} 2*3 + 1*(-1) \end{vmatrix} = \begin{vmatrix} 5 \end{vmatrix}$$

$$\begin{vmatrix} 0 & 3 \end{vmatrix} \begin{vmatrix} -1 \end{vmatrix} \begin{vmatrix} 0*3 + 3*(-1) \end{vmatrix} \begin{vmatrix} -3 \end{vmatrix}$$

### ■ Machine Learning Connection

Matrix multiplication is THE core computation in ML:

- **Forward pass in neural networks:**  $\text{output} = W * \text{input} + \text{bias}$  (matrix-vector multiply!)
- **Batch processing:** Multiply weight matrix by entire batch matrix at once
- **Attention mechanism (Transformers):**  $Q * K^T$  (matrix multiplication of query and key matrices)

```
Python: A = np.array([[1,2],[3,4]])
B = np.array([[5,6],[7,8]])
C = A @ B # or np.dot(A, B)
print(C) # [[19 22] [43 50]]
```

### Practice Problems

**P4.1** Calculate the dot product:  $[2, 3, 1] \cdot [4, -1, 5]$

**P4.2** Can you multiply a  $(3 \times 2)$  matrix by a  $(3 \times 2)$  matrix? Why or why not?

**P4.3** Multiply:  $|1 0| \times |5|$

$$|0 1| |3|$$

**P4.4** Multiply:  $|2 0| \times |1 3|$

$$|0 3| |2 4|$$

**P4.5** If A is  $5 \times 3$  and B is  $3 \times 7$ , what size is AB?

### Answers

**A4.1**  $2*4 + 3*(-1) + 1*5 = 8 - 3 + 5 = 10$

**A4.2** No! Columns of first (2) must equal rows of second (3). 2 is not 3.

**A4.3**  $[5, 3]$  - The identity matrix leaves vectors unchanged!

**A4.4**  $|2*1+0*2 2*3+0*4| = |2 6|$

$$|0*1+3*2 0*3+3*4| |6 12|$$

**A4.5**  $5 \times 7$

## Chapter 5

# Linear Transformations

MEDIUM

## 5.1 What is a Linear Transformation?

A **linear transformation** is a function that takes a vector as input and produces a vector as output, while preserving the operations of addition and scalar multiplication. The key insight: **every linear transformation can be represented as matrix multiplication!**

### Definition: Linear Transformation

A function  $T$  is linear if: (1)  $T(u + v) = T(u) + T(v)$ , and (2)  $T(cv) = cT(v)$  for any scalar  $c$ . In simple terms, it doesn't matter whether you add/scale before or after applying the transformation.

## 5.2 Common 2D Transformations

Transformation	Matrix	What it does
Scaling	$\begin{vmatrix} sx & 0 \\ 0 & sy \end{vmatrix}$	Stretches x by sx, y by sy
Rotation (90 deg)	$\begin{vmatrix} 0 & -1 \\ 1 & 0 \end{vmatrix}$	Rotates counterclockwise 90 degrees
Reflection (y-axis)	$\begin{vmatrix} -1 & 0 \\ 0 & 1 \end{vmatrix}$	Flips across the y-axis
Shear	$\begin{vmatrix} 1 & k \\ 0 & 1 \end{vmatrix}$	Tilts shapes sideways by factor k

### ■ Key Insight

**The Big Idea:** Every matrix IS a transformation. When you multiply a matrix by a vector, you are transforming that vector. The columns of the matrix tell you where the basis vectors land after the transformation.

For  $\begin{vmatrix} a & b \\ c & d \end{vmatrix}$  : the first column  $[a,c]$  is where  $[1,0]$  goes.

$[c,d]$  the second column  $[b,d]$  is where  $[0,1]$  goes.

## 5.3 Composition of Transformations

Applying transformation A then transformation B is the same as applying the single transformation BA (note: B comes first in the product because it acts last). This is why matrix multiplication is not commutative - rotating then scaling is different from scaling then rotating!

### ■ Machine Learning Connection

Neural networks are sequences of linear transformations (with non-linear activations in between):

- **Each layer:** applies a linear transformation (matrix multiply) then a non-linear activation
- **Deep learning:** composing many transformations = multiplying many matrices
- **Convolutional layers:** special linear transformations that detect patterns in images
- **Data augmentation:** rotation, scaling, and shearing matrices transform training images

### ↳ Practice Problems

**P5.1** Apply the scaling matrix  $\begin{vmatrix} 2 & 0 \\ 0 & 2 \end{vmatrix}$  to the vector  $[3, 1]$ . What happens?

$$\begin{vmatrix} 0 & 2 \\ 1 & 0 \end{vmatrix}$$

**P5.2** Apply the rotation matrix  $\begin{vmatrix} 0 & -1 \\ 1 & 0 \end{vmatrix}$  to  $[1, 0]$ . Where does it go?

$$\begin{vmatrix} 1 & 0 \\ 0 & 1 \end{vmatrix}$$

**P5.3** What happens when you apply the identity matrix to any vector?

**P5.4** If layer 1 has weight matrix W1 (3x5) and layer 2 has W2 (2x3), what is the size of an input vector, and what is the size of the final output?

## ■ Answers

**A5.1**  $|2 \ 0|^*|3| = |6|$  - The vector doubled in length (scaled by 2).

$|0 \ 2| \ |1| \ |2|$

**A5.2**  $|0 \ -1|^*|1| = |0|$  - The vector rotated 90 degrees counterclockwise!

$|1 \ 0| \ |0| \ |1|$

**A5.3** Nothing! The identity transformation leaves every vector unchanged.  $I^*v = v$ .

**A5.4** Input: 5-dimensional ( $5 \times 1$ ). After W1: 3-dimensional. After W2: 2-dimensional.

## Chapter 6

# Systems of Linear Equations

MEDIUM

## 6.1 What is a System of Equations?

A system of linear equations is a set of equations that must all be true simultaneously. Linear algebra gives us tools to solve these efficiently.

**Example:** Find  $x$  and  $y$  such that:

$$\begin{aligned} 2x + 3y &= 8 \\ x - y &= 1 \end{aligned}$$

This can be written as the matrix equation  $\mathbf{Ax} = \mathbf{b}$ :

$$\begin{array}{|c c|} \hline 2 & 3 \\ \hline \end{array} * \begin{array}{|c|} \hline x \\ \hline \end{array} = \begin{array}{|c|} \hline 8 \\ \hline \end{array}$$
  
$$\begin{array}{|c c|} \hline 1 & -1 \\ \hline \end{array} \begin{array}{|c|} \hline y \\ \hline \end{array} \quad | 1 |$$

## 6.2 Gaussian Elimination (Row Reduction)

The systematic way to solve systems. Write the **augmented matrix** and use row operations to reduce it to a simple form.

**Step-by-step for our example:**

**Step 1:** Write augmented matrix  $[A|b]$ :

$$\begin{array}{|c c|c|} \hline 2 & 3 & 8 \\ \hline 1 & -1 & 1 \\ \hline \end{array}$$

**Step 2:**  $R_2 = R_2 - (1/2)*R_1$  (eliminate  $x$  from row 2):

$$\begin{array}{|c c|c|} \hline 2 & 3 & 8 \\ \hline 0 & -5/2 & -3 \\ \hline \end{array}$$

**Step 3:** Back-substitute. From row 2:  $-5/2 * y = -3$ , so  $y = 6/5 = 1.2$

From row 1:  $2x + 3(1.2) = 8$ , so  $2x = 4.4$ ,  $x = 2.2$

## 6.3 Three Possible Outcomes

One Solution	No Solution	Infinite Solutions
Lines cross at exactly one point. The system is <b>consistent</b> and the matrix is <b>invertible</b> .	Lines are parallel (never cross). The equations contradict each other.	Lines overlap completely. There are more unknowns than constraints.

## 6.4 The Matrix Inverse

If  $A$  is an invertible (square) matrix, then  $A^{-1}$  exists such that  $A * A^{-1} = I$ . This lets us solve  $Ax = b$  directly:

$$Ax = b \rightarrow x = A^{-1} * b$$

2x2 Inverse Formula: For  $A = |a \ b| / |c \ d|$ :

$$A^{-1} = (1/(ad-bc)) * \begin{vmatrix} d & -b \\ -c & a \end{vmatrix}$$

The value  $(ad - bc)$  is the **determinant**. If it equals 0, the inverse does not exist!

### ■ Machine Learning Connection

Solving systems is central to ML optimization:

- **Linear regression (Normal Equation):**  $w = (X^T X)^{-1} X^T y$  solves for optimal weights directly
- **Least squares:** When no exact solution exists (overdetermined system), we find the best approximate solution
- **Regularization:** Ridge regression modifies the system to  $(X^T X + \lambda I)^{-1} X^T y$  for stability

Python: `w = np.linalg.solve(X.T @ X, X.T @ y) # Efficient!`

# Better than: `w = np.linalg.inv(X.T @ X) @ X.T @ y`

### Practice Problems

**P6.1** Write as matrix equation  $Ax = b$ :  $3x + y = 7$ ,  $2x - y = 3$

**P6.2** Find the inverse of  $|2 \ 1|$  using the 2x2 formula.

$|1 \ 1|$

**P6.3** Solve using the inverse you found:  $2x + y = 5$ ,  $x + y = 3$

**P6.4** Does  $|1 \ 2|$  have an inverse? (Hint: check the determinant)

$|2 \ 4|$

## ■ Answers

**A6.1**  $|3 \ 1| * |x| = |7|$

$|2 \ -1| |y| |3|$

**A6.2**  $\det = 2*1 - 1*1 = 1$ . Inverse =  $| 1 \ -1|$

$|-1 \ 2|$

**A6.3**  $x = | 1 \ -1| * |5| = |5-3| = |2|$ . So  $x=2, y=1$ .

$|-1 \ 2| |3| |-5+6| |1|$

**A6.4**  $\det = 1*4 - 2*2 = 0$ . No inverse exists! (rows are proportional)

## Chapter 7

# Determinants

ADVANCED

## 7.1 What is a Determinant?

The **determinant** is a single number computed from a square matrix that tells you important things about the matrix and the transformation it represents.

### ■ Key Insight

**Physical intuition:** The determinant measures how much a transformation **scales area** (in 2D) or **volume** (in 3D).

- $\det = 2$  means areas double
- $\det = 0.5$  means areas halve
- $\det = 0$  means everything collapses to a lower dimension (line or point)
- $\det < 0$  means the orientation flips (like looking in a mirror)

## 7.2 Computing Determinants

**2x2 Determinant:**

$$\det \begin{vmatrix} a & b \\ c & d \end{vmatrix} = ad - bc$$

$$\det \begin{vmatrix} 3 & 1 \\ 2 & 4 \end{vmatrix} = 3*4 - 1*2 = 10$$

**3x3 Determinant** (cofactor expansion along first row):

$$\begin{aligned} \det \begin{vmatrix} a & b & c \\ d & e & f \\ g & h & i \end{vmatrix} &= a(ei-fh) - b(di-fg) + c(dh-eg) \end{aligned}$$

## 7.3 Key Properties

- $\det(A) = 0$  means ~~A is singular~~ (not invertible) - the transformation collapses dimensions
- $\det(AB) = \det(A) * \det(B)$
- $\det(A^T) = \det(A)$
- $\det(cA) = c^n * \det(A)$  for an  $n \times n$  matrix
- Swapping two rows negates the determinant

### ■ Machine Learning Connection

Determinants help diagnose problems in ML:

- **Feature redundancy:** If  $\det(X^T X)$  is near zero, your features are highly correlated (multicollinearity)
- **Model stability:** A near-zero determinant means small data changes cause huge parameter changes
- **Gaussian distributions:** The multivariate normal distribution formula includes  $\det(\Sigma)$  for the covariance matrix
- **Volume of data:** The determinant of the covariance matrix measures how spread out your data is

### ☞ Practice Problems

**P7.1** Calculate:  $\det |5\ 3|$

$$|2\ 4|$$

**P7.2** Calculate:  $\det |1\ 0|$

$$|0\ 1| \text{ (identity matrix)}$$

**P7.3** If  $\det(A) = 3$  and  $\det(B) = 4$ , what is  $\det(AB)$ ?

**P7.4** A matrix has  $\det = 0$ . Can you solve  $Ax = b$  for a unique  $x$ ?

### ■ Answers

**A7.1**  $5*4 - 3*2 = 20 - 6 = 14$

**A7.2**  $1*1 - 0*0 = 1$  (the identity always has  $\det = 1$ )

**A7.3**  $\det(AB) = \det(A)*\det(B) = 3*4 = 12$

**A7.4** No!  $\det = 0$  means the matrix is singular (not invertible). No unique solution.

## Chapter 8

# Eigenvalues & Eigenvectors

ADVANCED

## 8.1 The Big Idea

When you apply a matrix transformation to most vectors, both their direction AND length change. But some special vectors only get **stretched or shrunk** - their direction stays the same (or flips 180 degrees). These are **eigenvectors**, and the stretch factor is the **eigenvalue**.

### Definition: Eigenvalue & Eigenvector

For a square matrix A, if  $Av = \lambda v$  (where v is not zero), then v is an **eigenvector** and  $\lambda$  is the corresponding **eigenvalue**. The transformation A simply scales v by  $\lambda$ .

**Example:** Let  $A = \begin{vmatrix} 2 & 1 \end{vmatrix}$  and  $v = \begin{vmatrix} 1 \end{vmatrix}$

$$\begin{vmatrix} 0 & 3 \end{vmatrix} \begin{vmatrix} 1 \end{vmatrix}$$

$$A \cdot v = \begin{vmatrix} 2 & 1 \end{vmatrix} \cdot \begin{vmatrix} 1 \end{vmatrix} = \begin{vmatrix} 3 \end{vmatrix} = 3 \cdot \begin{vmatrix} 1 \end{vmatrix}$$
$$\begin{vmatrix} 0 & 3 \end{vmatrix} \begin{vmatrix} 1 \end{vmatrix} \begin{vmatrix} 3 \end{vmatrix} \begin{vmatrix} 1 \end{vmatrix}$$

So  $[1, 1]$  is an eigenvector with eigenvalue 3. The matrix just tripled it!

## 8.2 How to Find Eigenvalues

Solve the **characteristic equation**:  $\det(A - \lambda I) = 0$

**Example:**  $A = \begin{vmatrix} 4 & 2 \end{vmatrix}$

$$\begin{vmatrix} 1 & 3 \end{vmatrix}$$

$$A - \lambda I = \begin{vmatrix} 4-\lambda & 2 \\ 1 & 3-\lambda \end{vmatrix}$$
$$\det = (4-\lambda)(3-\lambda) - 2 \cdot 1 = 0$$
$$\lambda^2 - 7\lambda + 10 = 0$$
$$(\lambda - 5)(\lambda - 2) = 0$$

---

```
lambda = 5 or lambda = 2
```

## 8.3 Key Properties

- An  $n \times n$  matrix has  $n$  eigenvalues (counting repeats, possibly complex)
- The sum of eigenvalues = trace of  $A$  (sum of diagonal elements)
- The product of eigenvalues = determinant of  $A$
- If any eigenvalue is 0, the matrix is singular (not invertible)
- Symmetric matrices always have real eigenvalues

### ■ Key Insight

**Why eigenvalues matter:** They reveal the "natural axes" of a transformation. Instead of a complex matrix operation, you can understand it as simple stretching along eigenvector directions. This is called **eigendecomposition**:  $A = V * D * V^{-1}$ , where  $D$  is a diagonal matrix of eigenvalues.

### ■ Machine Learning Connection

Eigenvalues and eigenvectors are critical in ML:

- **PCA (Principal Component Analysis):** Eigenvectors of the covariance matrix are the principal components. Eigenvalues tell you how much variance each component captures.
- **Google PageRank:** Web page importance is computed as the dominant eigenvector of the link matrix.
- **Spectral clustering:** Uses eigenvectors of the graph Laplacian to cluster data.
- **Stability analysis:** Eigenvalues of the Hessian matrix tell you if you're at a local minimum.

Python: `eigenvalues, eigenvectors = np.linalg.eig(A)`

### ➲ Practice Problems

**P8.1 Verify:** Is  $[1, 0]$  an eigenvector of  $\begin{vmatrix} 3 & 0 \\ 0 & 5 \end{vmatrix}$ ? If so, what is the eigenvalue?

$\begin{vmatrix} 0 & 5 \\ 0 & 1 \end{vmatrix}$

**P8.2 Find eigenvalues of**  $\begin{vmatrix} 2 & 0 \\ 0 & 0 \end{vmatrix}$ . (Hint: it's a diagonal matrix!)

$\begin{vmatrix} 0 & 7 \\ 0 & 0 \end{vmatrix}$

**P8.3 If eigenvalues of A are 3 and 5, what is  $\det(A)$ ?**

**P8.4 If eigenvalues of A are 3 and 5, what is the trace of A?**

## ■ Answers

**A8.1**  $|3 \ 0| * |1| = |3| = 3 * |1|$ . Yes! Eigenvalue = 3.

$|0 \ 5| \ |0| \ |0| \ |0|$

**A8.2** For diagonal matrices, eigenvalues ARE the diagonal entries: 2 and 7.

**A8.3**  $\det(A) = \text{product of eigenvalues} = 3 * 5 = 15$

**A8.4**  $\text{trace}(A) = \text{sum of eigenvalues} = 3 + 5 = 8$

## Chapter 9

# Norms, Dot Products & Distances

ADVANCED

## 9.1 Vector Norms (Length/Magnitude)

A **norm** measures the size (length) of a vector. There are several types:

Norm	Formula	Intuition	Example [3, -4]
L1 (Manhattan)	$\ v\ _1 =  v_1  +  v_2  + \dots$	Walk along grid lines	$3 + 4 = 7$
L2 (Euclidean)	$\ v\ _2 = \sqrt{v_1^2 + v_2^2 + \dots}$	Straight-line distance	$\sqrt{9+16} = 5$
L-infinity (Max)	$\ v\ _{\infty} = \max( v_i )$	Largest component	$\max(3,4) = 4$

## 9.2 Dot Product (Inner Product) Revisited

The dot product has a geometric interpretation:

$$a \cdot b = \|a\| * \|b\| * \cos(\theta)$$

where theta is the angle between the two vectors. This gives us:

- If  $a \cdot b = 0$ , the vectors are **orthogonal** (perpendicular, 90 degrees apart)
- If  $a \cdot b > 0$ , the vectors point in **similar** directions (angle < 90 degrees)
- If  $a \cdot b < 0$ , the vectors point in **opposite** directions (angle > 90 degrees)

## 9.3 Cosine Similarity

Cosine similarity normalizes the dot product to measure only the **angle** between vectors, ignoring their magnitudes:

$$\text{cos\_sim}(a, b) = (a \cdot b) / (\|a\| * \|b\|)$$

Result is always between -1 (opposite) and 1 (identical direction). This is one of the most used metrics in ML!

## 9.4 Orthogonality and Projection

The **projection** of vector  $a$  onto vector  $b$  gives you the component of  $a$  in the direction of  $b$ :

$$\text{proj}_b(a) = ((a \cdot b) / (b \cdot b)) * b$$

This is fundamental to least-squares regression and many optimization algorithms.

### ■ Machine Learning Connection

Norms and distances are everywhere in ML:

- **Loss functions:** L2 norm = MSE loss, L1 norm = MAE loss
- **Regularization:** L1 (Lasso) makes weights sparse; L2 (Ridge) makes weights small
- **KNN algorithm:** Uses Euclidean distance to find nearest neighbors
- **Cosine similarity:** Used in recommendation systems and NLP to compare documents/embeddings
- **Unit vectors:** Normalizing vectors ( $v / \|v\|$ ) is used in many algorithms (softmax, normalization layers)

```
Python: norm = np.linalg.norm(v) # L2 norm  
cos_sim = np.dot(a, b) / (np.linalg.norm(a) * np.linalg.norm(b))
```

### ☞ Practice Problems

- P9.1** Find the L2 norm of [3, 4]
- P9.2** Find the L1 norm of [3, -4, 2]
- P9.3** Are [1, 2] and [-2, 1] orthogonal? (Check their dot product)
- P9.4** Find the cosine similarity between [1, 0] and [1, 1]
- P9.5** Normalize [3, 4] to a unit vector (length 1)

### ■ Answers

- A9.1**  $\sqrt{9 + 16} = \sqrt{25} = 5$
- A9.2**  $|3| + |-4| + |2| = 3 + 4 + 2 = 9$
- A9.3**  $[1,2] \cdot [-2,1] = 1*(-2) + 2*1 = -2 + 2 = 0$ . Yes, orthogonal!
- A9.4**  $\text{dot} = 1*1 + 0*1 = 1$ .  $\|a\| = 1$ ,  $\|b\| = \sqrt{2}$ .  $\text{cos\_sim} = 1/\sqrt{2} = 0.707$
- A9.5**  $\|v\| = 5$ . Unit vector =  $[3/5, 4/5] = [0.6, 0.8]$

## Chapter 10

# Putting It All Together: ML Applications

ADVANCED

Now let's see how all these concepts come together in real machine learning applications.

## 10.1 Linear Regression: The Complete Picture

Linear regression finds weights  $w$  that minimize the error between predictions and actual values.

**Setup:** You have  $m$  data points with  $n$  features. Your data matrix  $X$  is  $(m \times n)$ , labels  $y$  is  $(m \times 1)$ , and weights  $w$  is  $(n \times 1)$ .

Predictions:  $\hat{y} = X * w$

Error:  $e = y - \hat{y}$

Loss:  $L = \|e\|_2^2 = (y - Xw)^T(y - Xw)$

**Solution (Normal Equation):**

$$w = (X^T X)^{-1} X^T y$$

**Linear algebra concepts used:** Matrix multiplication, transpose, inverse, systems of equations.

## 10.2 PCA: Dimensionality Reduction

Principal Component Analysis reduces the number of features while keeping the most important information.

Step	Action	Linear Algebra Used
1	Center the data (subtract mean)	Vector subtraction
2	Compute covariance matrix $C = (1/m) X^T X$	Matrix multiplication, transpose
3	Find eigenvalues and eigenvectors of $C$	Eigendecomposition
4	Sort eigenvectors by eigenvalue (largest first)	Eigenvalues = variance captured

5	Keep top k eigenvectors as new basis	Basis, span, dimension
6	Project data onto new basis: $X_{\text{new}} = X * V_k$	Matrix multiply, projection

## 10.3 Neural Networks: Linear Algebra in Action

A neural network is essentially a sequence of matrix operations with non-linear activations in between.

### Single layer forward pass:

$$\begin{aligned} z &= W * x + b \text{ (linear transformation)} \\ a &= \text{activation}(z) \text{ (non-linear function)} \end{aligned}$$

Where  $W$  is the weight matrix,  $x$  is the input vector,  $b$  is the bias vector.

### Full network:

Input (784-dim) -->  $W_1$  (256x784) --> ReLU -->  $W_2$  (128x256) --> ReLU -->  $W_3$  (10x128) --> Softmax --> Output (10-dim)

### Linear algebra concepts in backpropagation:

- Gradient computation uses the **transpose** of weight matrices
- Weight updates use **scalar multiplication** (learning rate) and **vector subtraction**
- Batch processing uses **matrix multiplication** to process many samples at once

## 10.4 Complete Concept Map

ML Concept	Linear Algebra Foundation
Data representation	Vectors and matrices
Linear regression	Matrix inverse, systems of equations, projection
Neural network layers	Matrix multiplication, linear transformations
PCA / Dimensionality reduction	Eigenvalues, eigenvectors, projection, basis change
Gradient descent	Vector operations, norms, scalar multiplication
Regularization (L1/L2)	Vector norms (L1, L2)

Similarity / distance metrics	Dot product, cosine similarity, norms
SVD / Matrix factorization	Eigendecomposition, orthogonality
Attention (Transformers)	Matrix multiplication, dot product, softmax
Convolutions (CNNs)	Specialized linear transformations

## 10.5 NumPy Quick Reference

Operation	Python Code
Create vector	<code>v = np.array([1, 2, 3])</code>
Create matrix	<code>A = np.array([[1,2],[3,4]])</code>
Matrix multiply	<code>C = A @ B # or np.dot(A, B)</code>
Transpose	<code>A.T # or np.transpose(A)</code>
Inverse	<code>np.linalg.inv(A)</code>
Determinant	<code>np.linalg.det(A)</code>
Eigenvalues	<code>vals, vecs = np.linalg.eig(A)</code>
Solve Ax=b	<code>x = np.linalg.solve(A, b)</code>
L2 norm	<code>np.linalg.norm(v)</code>
Dot product	<code>np.dot(a, b) # or a @ b</code>
Identity matrix	<code>np.eye(n) # n x n identity</code>
SVD	<code>U, S, Vt = np.linalg.svd(A)</code>

### Practice Problems

**P10.1** You have 100 data points with 5 features. What size is your data matrix X?

**P10.2** In linear regression with the matrix above, what size is the weight vector w?

**P10.3** What is  $X^T X$  for the above? What size is it?

**P10.4** A neural network layer has input size 256 and output size 128. What size is the weight matrix?

**P10.5** After PCA, you reduce 50 features to 10 principal components. How much smaller is your data?

## ■ Answers

**A10.1**  $100 \times 5$  (100 rows, 5 columns)

**A10.2**  $5 \times 1$  (one weight per feature)

**A10.3**  $(5 \times 100) * (100 \times 5) = 5 \times 5$  matrix (the Gram matrix)

**A10.4**  $128 \times 256$  (output\_dim x input\_dim)

**A10.5** From 50 columns to 10 columns = 5x reduction in feature dimensionality

# Your Learning Path Forward

Congratulations on working through this guide! Here's how to continue building your linear algebra skills for machine learning:

Stage	Focus	Resources
1. Practice	Redo all practice problems without looking at answers. Implement each operation in NumPy.	This guide + Python
2. Deepen	3Blue1Brown "Essence of Linear Algebra" (YouTube). Best visual explanations available.	YouTube (free)
3. Apply	Implement linear regression, PCA, and a simple neural network from scratch using only NumPy.	Python + NumPy
4. Expand	Learn SVD (Singular Value Decomposition), matrix factorization, and sparse matrices.	Gilbert Strang textbook
5. Master	Study the math behind Transformers, attention mechanisms, and optimization algorithms.	Research papers

## ■ Key Insight

**Remember:** Linear algebra is a skill, not just knowledge. You build it by doing, not just reading. Write code, solve problems by hand, and most importantly - when you encounter a new ML concept, ask yourself: "What linear algebra is happening here?"

Every matrix operation you understand makes the next ML paper or algorithm easier to grasp. Keep going!

**Linear Algebra + Machine Learning = Your Superpower**